

YOUR WEEKLY MICROSOFT SENTINEL BRIEFING

Sentinel Weekly

NO. 03 · AUGUST 6 TO 12, 2026

*The playbook generator drops its
Security Copilot requirement.*

Subscribe to Sentinel Weekly to get every edition the moment it drops.

Edition 03 covers Thursday, August 6 to Wednesday, August 12, 2026. One note on sourcing: Microsoft's monthly roundups for Sentinel and Defender landed inside this same window, and they repeat a lot of what Editions 01 and 02 already covered. Only the genuinely new items from those pages made it in here.

On the radar this week

- **The playbook generator no longer needs Security Copilot** (*Microsoft Learn*) No separate Security Copilot license, no Security Compute Units, no extra cost. The single biggest reason most teams never switched it on is gone.
- **UEBA picks up FortiGate, Check Point, Zscaler and AWS GuardDuty** (*Microsoft Learn, August 2026*) More than 40 new FortiGate behaviors from CommonSecurityLog, ten new anomaly rules across three firewall and proxy vendors, and identity linked findings from the AWSGuardDuty table.
- **Defender starts detecting threats against Agent 365 agents** (*Microsoft Learn*) Runtime signals from agent interactions, tool usage and execution patterns now raise alerts in Defender XDR, and those alerts are queryable in advanced hunting.
- **Sentinel retires in Azure operated by 21Vianet on August 18** (*Microsoft Learn*) Every Sentinel feature goes away in that region in five days. New subscription onboarding is already closed.

In this edition

The playbook generator loses its paywall and gains a long list of limits worth reading. UEBA quietly doubles its firewall coverage. Defender starts alerting on AI agents while the blocking half goes GA ahead of it. Plus a five day deadline, and a query for deciding what to automate first.

Shipped and live now

Generate playbooks using AI in Microsoft Sentinel (*Microsoft Learn*)

The playbook generator is no longer gated behind Security Copilot. Microsoft confirmed on August 6 that it needs neither a separate Security Copilot license nor Security Compute Units. It runs inside the Defender portal, in an embedded VS Code environment driven by Cline, and it writes Python rather than Logic Apps.

The workflow is Plan mode, then Act mode. You describe the response you want in natural language, the generator returns a plan and a flow diagram, you approve it, and then it produces the code, the documentation and a test run against a real alert ID that you nominate. Permissions are unified RBAC: Automation Playbooks (Read and Write) to generate and deploy, plus Microsoft Sentinel Contributor on the relevant workspace if you also want to author the automation rules. Allow up to two hours for role assignments to take effect.

Read the limits before you plan anything large around this. Alerts are the only supported input type, external libraries are not supported, playbooks cannot call other playbooks, and there is no automatic code validation, so somebody still has to read the Python. The ceilings are 100 playbooks and 500 integrations per tenant, 5,000 lines and a 10 minute runtime per playbook, and 8M AI tokens per tenant per day. New playbooks are created disabled, so nothing runs until you switch them on.

Two operational details belong in your runbook today. Enhanced Alert Trigger rules live in a separate automation rules table from your Standard Alert Trigger rules, and there is no automatic migration, so expect to run both side by side for a while. And run results are not written to the SentinelHealth table: you find them in the Activities tab of the incident that holds the alert. If you monitor playbook failures through SentinelHealth today, that monitoring will be completely blind to every generated playbook you deploy.

New data sources for UEBA behaviors and anomaly detection *(Microsoft Learn, August 2026)*

UEBA picked up four new vendors this month, all in preview. Fortinet FortiGate firewall events from the CommonSecurityLog table now feed the behaviors layer with more than 40 new behaviors covering administrative activity on the appliances: rapid system reconfigurations, configuration backups, certificate changes and disruption of security services. They map to T1685, T1070, T1078 and T1602.002.

On the anomaly side, ten new rules read Check Point, Fortinet FortiGate and Zscaler firewall, VPN and web proxy events, also from CommonSecurityLog, and compare each user and device both against its own history and against organization wide activity. They cover anomalous and failed VPN sign ins, unusual access to high risk web categories, bursts of security detections on a single potentially compromised device, and suspicious administrative changes. The mappings are T1078, T1133, T1110, T1071, T1562 and T1567, and events arrive enriched with source IP geolocation, ISP and threat intelligence context. Identity linked AWS GuardDuty findings from the AWSGuardDuty table joined in the same batch, mapped to T1078, T1078.004, T1110, T1087.004 and T1567.002.

If your firewall and proxy logs already land in CommonSecurityLog, this is close to free coverage. The catch is that behaviors only build for sources you have explicitly enabled for UEBA, so open your data connector configuration and confirm the relevant connectors are actually flagged before you go looking for results that will not be there.

Public Preview: Nested API Support Comes to Microsoft Sentinel CCF *(Sentinel Blog, August 5)*

A small change that unblocks a large category of connectors. Nested API support lets a single Codeless Connector Framework polling cycle span multiple dependent calls: the first call returns a list of records such as alert IDs or case references, and one or more follow up calls fetch the full detail for each record in that list. That two step shape is how a great many ISV security APIs work, and until now it pushed you out of the codeless framework and into a Logic App or a Function. If you have a connector project that stalled on exactly this pattern, it is worth reopening.

Agentic SOC watch

Detect and investigate threats to AI agents using Microsoft Defender *(Microsoft Learn, Preview)*

Defender now analyzes runtime signals from agent interactions, tool usage and execution patterns, and surfaces the results as alerts in Defender XDR. Coverage spans observability data from Copilot Studio, Microsoft Foundry, the Microsoft 365 Copilot Agent Builder and anything integrated through the Microsoft Agent 365 SDK. The alerts land in incidents and in advanced hunting, which means this is queryable rather than just another dashboard.

Protect AI agents in real time using Microsoft Defender *(Microsoft Learn, GA)*

The enforcement half shipped to general availability in the same batch. Defender evaluates tool invocations and responses against security policy and can allow or block interactions with Work IQ MCP and customer MCP tools onboarded to Agent 365. Note the ordering: blocking is GA while detection is still preview, so you can enforce policy before you can fully hunt across what it caught. Plan the rollout accordingly, and keep the first policies in audit where you can.

The practical question for most SOCs is not the technology, it is ownership. Agent alerts arrive in the same incident queue as everything else, and almost nobody has written the triage path for them yet. Decide now who picks up an alert that says a Copilot Studio agent invoked an unexpected tool, because the first one will not wait for your process.

Next door: Defender XDR and Entra

What's New in Microsoft Entra: August 2026 *(Entra Blog, August)*

Three items from the Entra roundup carry real SOC impact. Administrators can now auto accept single sign on permission prompts on managed Windows devices through a supported registry setting, delivered with the July 2026 security update for Windows 11 versions 24H2 and 25H2 as KB5101650. Expect the shape of your consent prompt telemetry to change once that rolls out, and check any detection that keys on prompt frequency.

Microsoft Entra Domain Services now keeps automatic backups of managed group policy objects with a restore path, which finally gives you a recovery option after a malicious or simply careless policy change. And Microsoft Authenticator added jailbreak and root detection: secure by default, no admin configuration required, blocking Entra credentials from being added or used on compromised iOS and Android devices.

The clock to watch

Microsoft Sentinel feature support for Azure commercial and other clouds *(Microsoft Learn)*

Five days. Every Microsoft Sentinel feature retires in the Azure operated by 21Vianet region on August 18, 2026, and onboarding of new subscriptions there closed some time ago. If you run or advise anyone in that region, the window for an orderly exit has effectively shut, and Microsoft's standing guidance is to work the impact assessment through your account representative.

Update: the containerized SAP data connector deadline from Edition 02 has not moved. Support ends on September 30, 2026, and all new deployments now get the agentless connector only, billed at the same price. Six weeks left on that one.

One query to steal

Now that the playbook generator costs nothing extra, the bottleneck moves from budget to judgement: what do you automate first? Repetition is the answer, and repetition is measurable.

```
// illustrative, validate in your workspace first
// Which alerts actually repeat often enough to be worth automating?
let lookback = 30d;
let minDays = 10;
let minCount = 20;
SecurityAlert
| where TimeGenerated > ago(lookback)
| where Status != "Dismissed"
| extend Ents = todynamic(Entities)
| summarize
    Volume = count(),
    ActiveDays = dcount(bin(TimeGenerated, 1d)),
    Severities = make_set(AlertSeverity, 5),
    Providers = make_set(ProviderName, 5),
    EntityTypes = make_set(tostring(Ents[0].Type), 5)
by AlertName
| where ActiveDays >= minDays and Volume >= minCount
| extend PerDay = round(Volume * 1.0 / ActiveDays, 1)
| project AlertName, Volume, ActiveDays, PerDay, Severities, Providers, EntityTypes
| sort by Volume desc
| take 25
```

Sort by volume, then read the ActiveDays column before you decide anything. An alert firing 400 times across three days is an incident and needs a human. An alert firing 200 times across 28 days is a habit, and habits are what the generator is good at. Start with the second kind, and take the EntityTypes column with you into the prompt: the generator needs to know what it is being handed.

That's Edition 3

What landed in your tenant this week that I missed, and what should Edition 04 dig into? Drop it in the comments. I read every one, and reader suggestions have shaped every edition so far.

[Never miss an edition: subscribe to Sentinel Weekly on LinkedIn](#)

[Every past edition is archived as a PDF on GitHub](#)

Curated by Phillipe Chaves de Lima. New every Thursday.