

YOUR WEEKLY MICROSOFT SENTINEL BRIEFING

Sentinel Weekly

EDITION **NO. 5** · AUGUST 20 TO 26, 2026

*Linux gets an audit mode before you switch enforcement on.
Plus Store apps in vuln management and a Conditional Access
clock.*

IN THIS EDITION

- Shipped and live now: an audit mode for Linux, and Store apps enter vulnerability management
- Agentic SOC watch: MDASH reaches Azure DevOps, agent runtime protection widens
- Next door: RPC auditing turns itself on, password protection reaches SaaS
- The clock to watch: baseline scopes enforcement, and SAP content that renamed itself
- One query to steal: sizing the blast radius before you enforce on Linux

Sentinel Weekly, Edition 5. [Subscribe on LinkedIn](#) to get every edition the moment it drops.

Edition 5 covers August 20 to 26, 2026. The Sentinel what's new page did not move this week, and Editions 3 and 4 already worked through everything it had for August, so the weight of this edition sits on the endpoint and identity side, where the Defender for Endpoint page was refreshed inside the window. Microsoft's monthly roundups repeat a lot of what you already read, so only the genuinely new parts made the cut, and one item from Edition 2 returns below as a labelled update rather than a fresh story.

On the radar this week

- **Antivirus audit mode for Linux** (*Microsoft Learn*) Real time detection and alerting with no quarantine and no process termination, so you can measure efficacy on sensitive workloads before you enforce. Preview from build 101.26062.0007.
- **Vulnerability assessment reaches Microsoft Store applications** (*Microsoft Learn*) Teams, Firefox, WhatsApp, Slack, Dropbox and more, with the vulnerable file path and a suggested query in the Software evidence area. Preview.
- **Conditional Access enforcement for baseline scopes** (*Microsoft Learn*) The rollout that started on June 15 is landing in tenants now. It can add MFA prompts to Azure CLI and Visual Studio Code sign ins that were previously waved through.
- **Three SAP solution releases, and a rule rename** (*GitHub, Azure Sentinel*) Agentless 1.1.12, BTP 3.1.1 and LogServ 3.0.5. The BTP rename and the HANA move to Syslog will both break content that is pinned to the old names.

In this edition

- Shipped and live now: an audit mode for Linux, and Store apps enter vulnerability management
- Agentic SOC watch: MDASH reaches Azure DevOps, agent runtime protection widens
- Next door: RPC auditing turns itself on, password protection reaches SaaS
- The clock to watch: baseline scopes enforcement, and SAP content that renamed itself
- One query to steal: sizing the blast radius before you enforce on Linux

Shipped and live now

Antivirus audit mode for Linux, and an offboarding API to match (*Microsoft Learn, Defender for Endpoint what's new, August 2026*)

Defender for Endpoint on Linux picked up an antivirus audit mode in preview, and it is the most practically useful thing to land this week. It gives you real time malware detection and alerting without quarantining files or terminating processes. That is exactly the mode you want before turning enforcement on across a fleet of database, build and application servers, where a false positive that quarantines a file or kills a process is an outage rather than an inconvenience.

You get the detection telemetry and the alerts, the workload keeps running, and you can measure both efficacy and performance impact before committing to enforcement. It needs Defender for Endpoint on Linux version 101.26062.0007 or later.

The same version takes offboarding API support for Linux to GA, so the Offboard machine API now works for Linux servers and decommissioning can be automated instead of done by hand. If your Linux estate churns,

those two are worth planning a single rollout around. Two macOS builds shipped in the same window as well, 101.26062.0011 and 101.26062.0012.

The trap to avoid: audit mode is only worth running if somebody actually reads the output. Give it a fixed window, two or three weeks, and an owner who has to produce a list of what would have been quarantined and why. Audit mode left running indefinitely is just an unprotected fleet with better logging.

Vulnerability assessment reaches Microsoft Store applications (Preview) *(Microsoft Learn, Defender for Endpoint what's new, August 2026)*

Defender Vulnerability Management can now assess vulnerabilities in Microsoft Store applications. The named list includes Microsoft Teams, Mozilla Firefox, WhatsApp, Slack, Dropbox, DuckDuckGo, Dell Command, HP Smart, HP Support Assistant, Dell SupportAssist for Home PCs and NVIDIA Control Panel. The Software evidence area on the software page gives you the file path for the vulnerable application, plus a suggested query showing the app, its version and its file location.

This closes a blind spot that was easy to miss. A Store installed browser and the same browser installed from an MSI were not the same thing to your inventory, and the Store side was usually the one nobody was patching, because nobody knew it was there. Expect your device level exposure numbers to move once this lights up, and expect at least one argument about whether the new findings are real or an artefact of better visibility. They are real. They were always there.

Agentic SOC watch

Update: MDASH picks up an Azure DevOps connector (Preview) *(Microsoft Learn, Exposure Management what's new)*

We covered the MDASH agentic code scanner and the MAI-Cyber-1-Flash model in Edition 2. The status has moved, so here is the update rather than the story again: the Azure DevOps connector is now in preview. You can connect Azure DevOps organizations from the Defender portal, onboard repositories, and trigger remote on demand agentic scans against them.

Until now that path ran through a GitHub connector or the Defender CLI. If your code lives in Azure DevOps, this is the point where the scanner becomes something you can actually pilot rather than something you read about. Scan results land in the Defender portal alongside the rest of your exposure data.

AI agent runtime protection widens its net *(Microsoft Learn, Defender for Endpoint what's new)*

Runtime protection for local AI agents on endpoints got two changes worth noting. Vendor supported agent event interfaces now work with the standard platform and engine update channels, so you no longer have to park machines on the Beta channel to get it. Agent native event inspection now covers Codex CLI and the GitHub Copilot app. And for agents that expose no vendor supported event interface at all, including OpenClaw and similar Node.js based Claw agents, Defender falls back to network inspection.

For detection engineers the interesting part is that fallback. Coverage is no longer strictly gated on a vendor shipping an interface, but it also means the fidelity of your agent alerts now varies by agent. Before you write detections on top of this, work out which of your agents are seen natively and which are only seen on the wire, because the two produce very different evidence.

Next door: Defender XDR and Entra

Defender for Identity turns on RPC auditing for you *(Microsoft Learn, Defender for Identity what's new)*

Enhanced RPC auditing is a prerequisite for several advanced Defender for Identity detections, and until now you had to apply a tag manually on each domain controller to switch it on. From sensor version 3.0.8 and later, Defender for Identity enables RPC auditing automatically when you upgrade.

If you have been carrying the Sensor v3.x RPC Audit Misconfigured health alert, this is the fix. Check the sensor version on your domain controllers, then check the health alert afterwards rather than assuming it cleared. Detection coverage that depends on a prerequisite nobody applied is the quietest kind of gap.

Password protection now reaches SaaS applications (Preview) *(Microsoft Learn, Defender for Identity what's new)*

The Password protection page now shows password risk from SaaS apps connected through Defender for Cloud Apps, on top of Active Directory, Microsoft Entra ID and Okta. SaaS apps that support SaaS Security Posture Management, Salesforce and ServiceNow among them, appear on the Password Hygiene and Password Policies tabs. Each SaaS app needs its own Defender for Cloud Apps app connector, so this is not free visibility, it is visibility you have to go and connect.

Worth doing anyway. Weak and reused passwords in a SaaS app that sits outside your identity provider are the ones that never show up in a Conditional Access report, and they are exactly where credential stuffing lands first.

The clock to watch

Improved enforcement for All resources policies with resource exclusions *(Microsoft Learn, Microsoft Entra ID)*

This is the one to read properly. Microsoft Entra ID is changing how Conditional Access policies that target All resources and carry one or more resource exclusions are enforced. Previously, baseline scopes were automatically excluded from enforcement whenever such an exclusion existed. Now they are evaluated as directory access and are subject to your policies.

Baseline scopes means the OIDC scopes email, offline_access, openid and profile, plus the baseline directory scopes User.Read, User.Read.All, User.ReadBasic.All, People.Read, People.Read.All, GroupMember.Read.All and Member.Read.Hidden.

The practical effect: a user signing in to the Visual Studio Code desktop client, which asks only for openid and profile, or to Azure CLI, which asks only for User.Read, can now be prompted for MFA or stopped by a device compliance control where they were waved through before. Applications that already request anything beyond the baseline are unaffected, which is most of them.

Rollout began on June 15, 2026 and runs progressively over several weeks, which puts late August squarely inside the window for a lot of tenants. You are only affected if you have All resources policies with resource exclusions and users signing in through apps that request nothing but baseline scopes. Microsoft documents a Graph query against auditLogs/signIns filtered on conditionalAccessAudiences that lists exactly those apps, so run it before enforcement finds them for you. Then either enable enforcement early in a test tenant, or use Customize behavior to keep the legacy behavior for named policies. Disabling enforcement tenant wide is possible and is not recommended.

SAP solution releases: 1.1.12, 3.1.1 and 3.0.5 *(GitHub, Azure Sentinel release notes)*

Three SAP solution releases landed. The agentless solution reached 1.1.12 with audit log performance improvements and a new force-sal-filesystem parameter. The BTP solution reached 3.1.1, adding an analytic rule for unaudited custom apps with login only activity, and renaming its analytic rules with an SAP BTP prefix. The LogServ solution reached 3.0.5, which lets customers on RISE with SAP activate existing ASIM based security content against standard SAP LogServ logs, and moves HANA detections from the custom log onto Syslog with updated connector coverage and fallback handling.

The rename is the part that bites. Anything keyed on analytic rule display names needs checking against the new prefix: automation rules, playbook conditions, ticket routing, workbooks and any reporting you built on

rule name. The HANA source change is the second one. If you have detections or visualizations pinned to the old custom log table, they will go quiet rather than throw an error, which is the failure mode nobody notices until an audit.

One query to steal

Audit mode is only worth running if you turn its output into a number you can defend. This one sizes the blast radius: what Defender would have quarantined or terminated on your Linux estate, grouped so you can see whether the answer is three stray test binaries or a build agent that would have been decapitated.

```
// illustrative, validate in your workspace first
// Pre enforcement impact check for Defender for Endpoint on Linux
// running in antivirus audit mode. Run it over the audit window.
let lookback = 14d;
let linuxDevices =
    DeviceInfo
    | where Timestamp > ago(lookback)
    | where OSPlatform startswith "Linux"
    | summarize arg_max(Timestamp, DeviceName, OSPlatform) by DeviceId
    | project DeviceId, DeviceName, OSPlatform;
let malwareAlerts =
    AlertInfo
    | where Timestamp > ago(lookback)
    | where ServiceSource == "Microsoft Defender for Endpoint"
    | where Category in ("Malware", "UnwantedSoftware", "SuspiciousActivity")
    | project AlertId, AlertTime = Timestamp, Title, Severity, Category;
AlertEvidence
| where Timestamp > ago(lookback)
| where isnotempty(DeviceId) and EntityType == "File"
| join kind=inner linuxDevices on DeviceId
| join kind=inner malwareAlerts on AlertId
| extend FilePath = strcat(tostring(AdditionalFields.FolderPath), "/", FileName)
| summarize
    Detections = count(),
    Devices = dcount(DeviceId),
    DeviceList = make_set(DeviceName, 20),
    Severities = make_set(Severity),
    FirstSeen = min(AlertTime),
    LastSeen = max(AlertTime)
    by Title, FileName, SHA256
| sort by Devices desc, Detections desc
```

Read the Devices column before the Detections column. One noisy file on one build agent is a tuning exercise; the same file on forty hosts is a deployment artefact somebody shipped on purpose, and enforcing on it will break a pipeline at 03:00. Confirm the column names against your own schema first, because AlertEvidence carries different fields depending on which products are onboarded, and add an exclusion list for the paths you already know are noisy before you show this to anyone who signs off on the rollout.

That's Edition 5

What should be in Edition 6? Drop the thing you are actually wrestling with in the comments, whether that is an audit mode rollout nobody will sign off on, a Conditional Access surprise that landed without warning, or an SAP connector that went quiet, and I will chase it down for next week.

Never miss an edition: [subscribe to Sentinel Weekly on LinkedIn](#)

PDF archive of every edition: github.com/PhillipeChaves/sentinel-weekly

Curated by Phillipe Chaves de Lima. New every Thursday.